

PROYECTO FINAL – ARQUITECTURA MEDALLION CON AZURE + DATABRICKS

Autor: César Fernando Campos Millán

Curso: Ingeniería de Datos e IA con Databricks

Fecha: Noviembre 2025

1. Introducción

El presente proyecto implementa un flujo completo de procesamiento de datos utilizando la Arquitectura Medallion sobre Azure Databricks y Azure Data Lake Storage (ADLS).

El objetivo principal es construir un pipeline moderno y escalable que permita:

- Ingestar datos crudos desde ADLS
- Transformarlos mediante PySpark
- Organizar los datos en capas (Bronze, Silver, Gold)
- Publicar tablas curadoradas en Delta Lake
- Automatizar el ETL con Databricks Jobs
- Crear dashboards analíticos basados en la capa Gold

2. Arquitectura del Proyecto

El pipeline sigue la arquitectura Medallion:

RAW → BRONZE → SILVER → GOLD → Dashboards

Servicios utilizados:

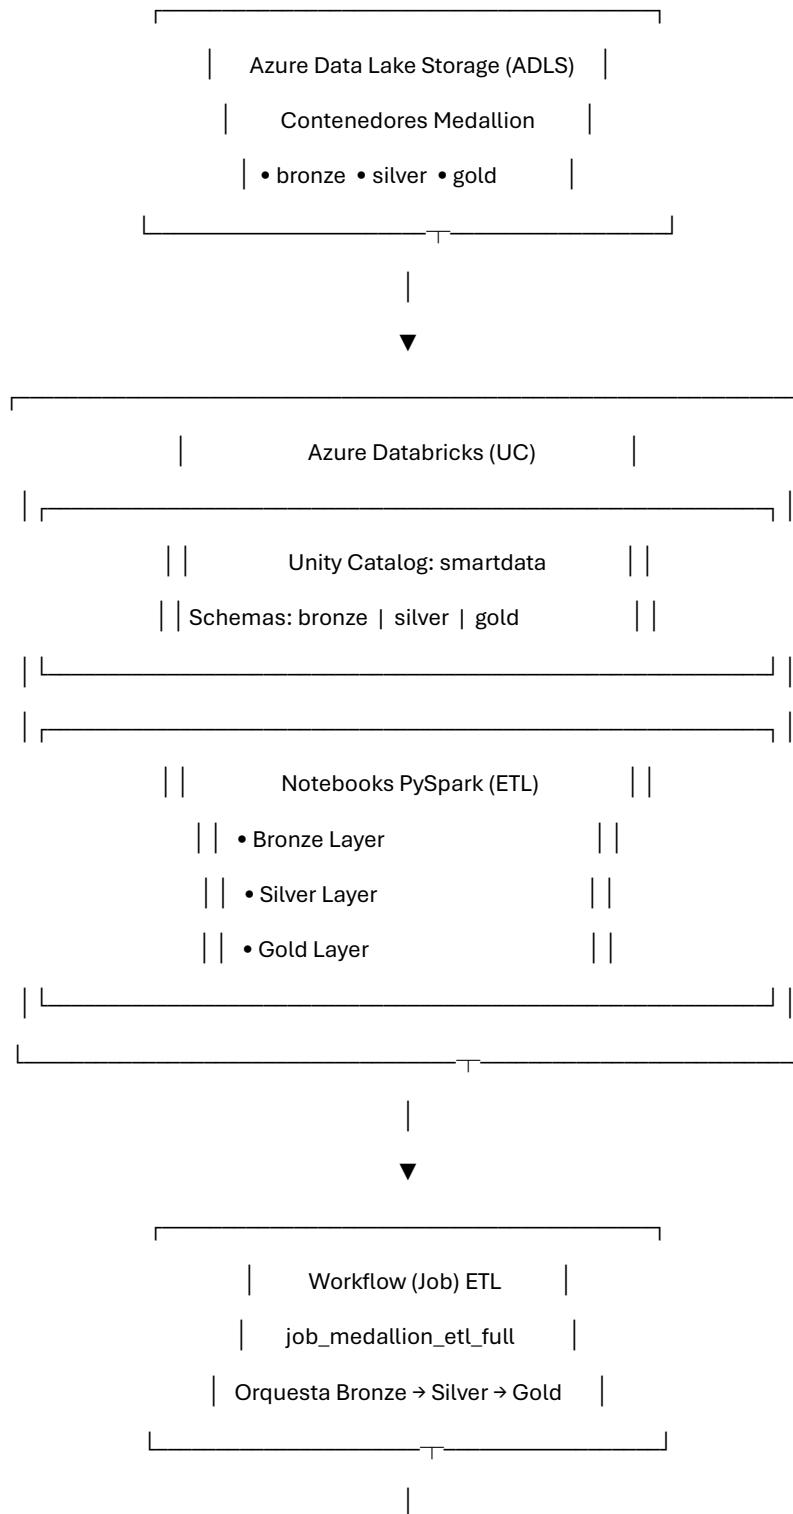
Azure Data Lake Storage Gen2 (ADLS)

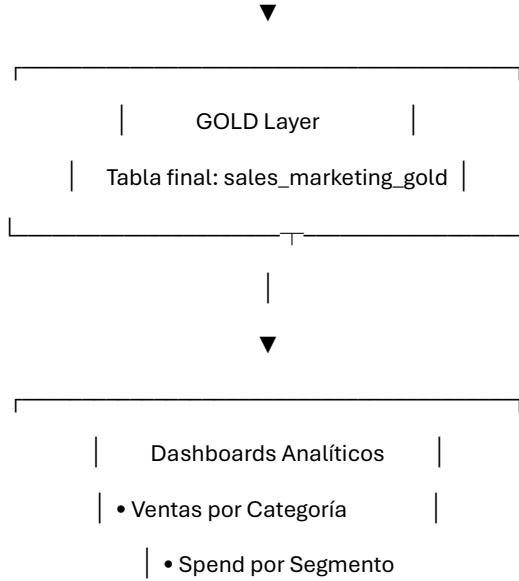
Azure Databricks (Unity Catalog)

Delta Lake

Databricks Jobs / Workflows

Databricks SQL Dashboards





3. Azure Data Lake – Capa de Almacenamiento

Se configuró un Storage Account con tres contenedores:

bronze

silver

gold

Cada contenedor representa una capa de la arquitectura Medallion.

The screenshot shows the Microsoft Azure Blob Storage interface. The left sidebar includes links for Overview, Activity log, Tags, Diagnose and solve problems, Access Control (IAM), Data migration, Events, Storage browser, Partner solutions, Resource visualizer, Data storage, Containers (which is selected), and File shares. The main area displays a table of containers with columns for Name, Last modified, Anonymous access level, and Lease state. The containers listed are Slogs, bronze, gold, metastore, and silver, all of which are private and available.

Name	Last modified	Anonymous access level	Lease state
Slogs	12/11/2025, 17:33:39	Private	Available
bronze	12/11/2025, 17:36:46	Private	Available
gold	12/11/2025, 17:37:46	Private	Available
metastore	12/11/2025, 22:33:32	Private	Available
silver	12/11/2025, 17:36:56	Private	Available

4. Capa BRONZE – Ingesta de Datos

En esta capa se almacenan los datos crudos provenientes de:

DBFS (marketing_campaign.csv)

ADLS (Ecommerce_Sales_Prediction_Dataset.csv)

Transformaciones realizadas:

Lectura en formato CSV

Inferencia automática de esquemas

Normalización mínima (renombrado de columnas cuando aplica)

Código ejemplo Bronze:

```
spark.sql("USE CATALOG smartdata")
spark.sql("USE SCHEMA bronze")
```

```
df_mkt_bronze = (
    spark.read.format("csv")
```

```
.option("header", "true")
.option("inferSchema", "true")
.option("delimiter", "\t")
.load("dbfs:/FileStore/marketing_campaign.csv")
)
df_mkt_bronze.write.format("delta").mode("overwrite") \
.saveAsTable("smartdata.bronze.marketing_raw")
```

Tablas generadas:

smartdata.bronze.marketing_raw
smartdata.bronze.ecommerce_raw

5. Capa SILVER – Limpieza y Estandarización

En esta etapa se realizan transformaciones para mejorar la calidad del dato.

Acciones principales:

Conversión de tipos: int, double, date

Corrección de formatos

Estandarización de columnas

Preparación para capa Gold

Código ejemplo Silver:

```
df_marketing_silver = (
    df_marketing
    .withColumn("Dt_Customer", F.to_date("Dt_Customer", "yyyy-MM-dd"))
    .withColumn("Income", F.col("Income").cast("int"))
)
```

```
df_marketing_silver.write.format("delta").mode("overwrite") \
```

```
.saveAsTable("smartdata.silver.marketing_campaign_silver")
```

6. Capa GOLD – Curación y Enriquecimiento

La capa GOLD contiene la versión “curada” y enriquecida del dato.

Acciones realizadas:

Cálculo de nuevas métricas:

Age

Total_Spend

Net_Price

Revenue

Generación de tabla final unificada para análisis:

```
smartdata.gold.sales_marketing_gold
```

Código ejemplo Gold:

```
df_gold = df_ecom_silver.crossJoin(df_mkt_silver)
df_gold = df_gold.withColumn("row_id", F.monotonically_increasing_id())

df_gold.write.format("delta").mode("overwrite") \
    .saveAsTable("smartdata.gold.sales_marketing_gold")
```

7. Unity Catalog – Gobernanza del Proyecto

Se organizó el catálogo:

Catálogo: smartdata

Schemas:

bronze

silver

gold

Tablas creadas:

marketing_raw

ecommerce_raw

marketing_campaign_silver

ecommerce_silver

marketing_campaign_gold

ecommerce_gold

sales_marketing_gold

The screenshot shows the Databricks Catalog interface. On the left, the sidebar includes options like Workspace, Recents, Catalog (which is selected), Jobs & Pipelines, Compute, Marketplace, SQL, SQL Editor, Queries, Dashboards, Genie, Alerts, Query History, and SQL Warehouses. The main area has tabs for Catalog, Delta Sharing, Clean Rooms, External Data, and Browse DBFS. A search bar at the top says "Search data, notebooks, recents, and more...". Below it, there's a "Type to search..." input field and a "Governance" dropdown. The Catalog section displays a tree view of the "smartdata" catalog, including "system", "main", "bronze", "gold", "information_schema", "silver", "Delta Shares Received", and "Legacy". To the right, a "Quick access" section shows a list of recent items: "smartdata" (Catalog, last viewed 1 day ago), "bronze" (Schema, last viewed 1 day ago), "gold" (Schema, last viewed 1 day ago), and "silver" (Schema, last viewed 1 day ago). There are also buttons for Recents, Favorites, and Catalogs, and a "Filter" search bar.

⚙️ 8. Databricks Workflow (Job)

Se implementó un job llamado:

📝 job_medallion_etl_full

Este job ejecuta automáticamente las tres capas del ETL.

Funcionalidades:

Automatiza el pipeline

Orquesta Bronze → Silver → Gold

Permite ejecución manual o programada

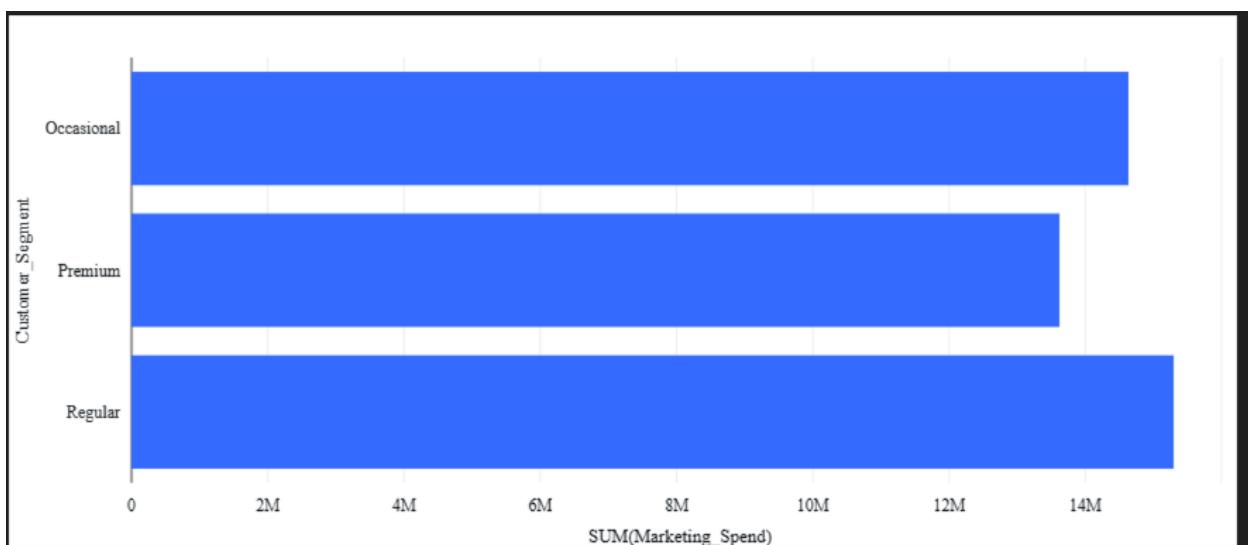
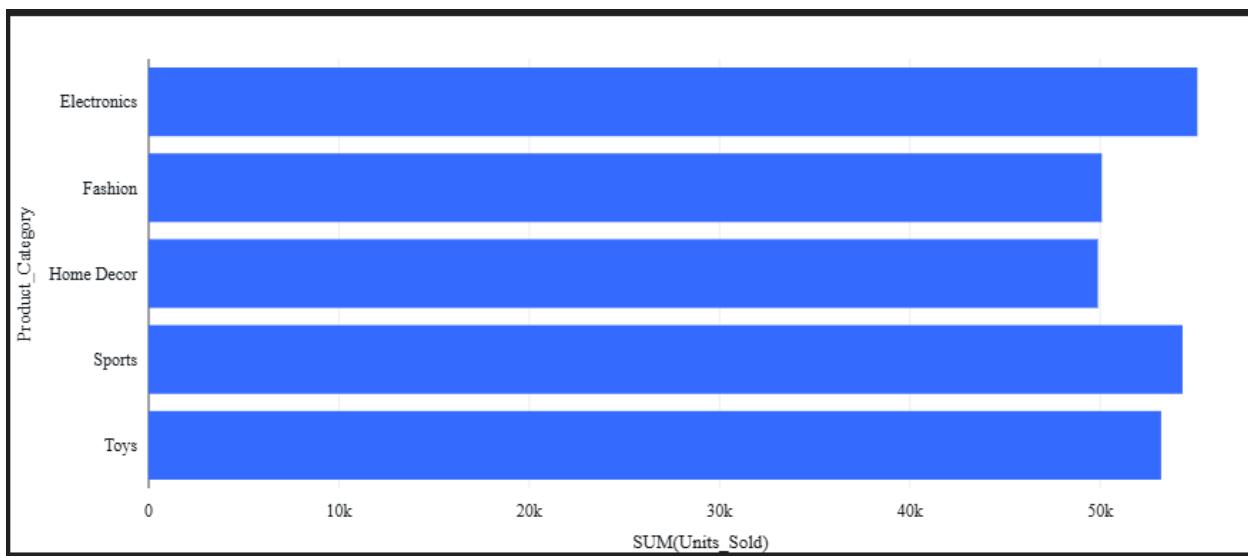
Conecta al cluster de Databricks

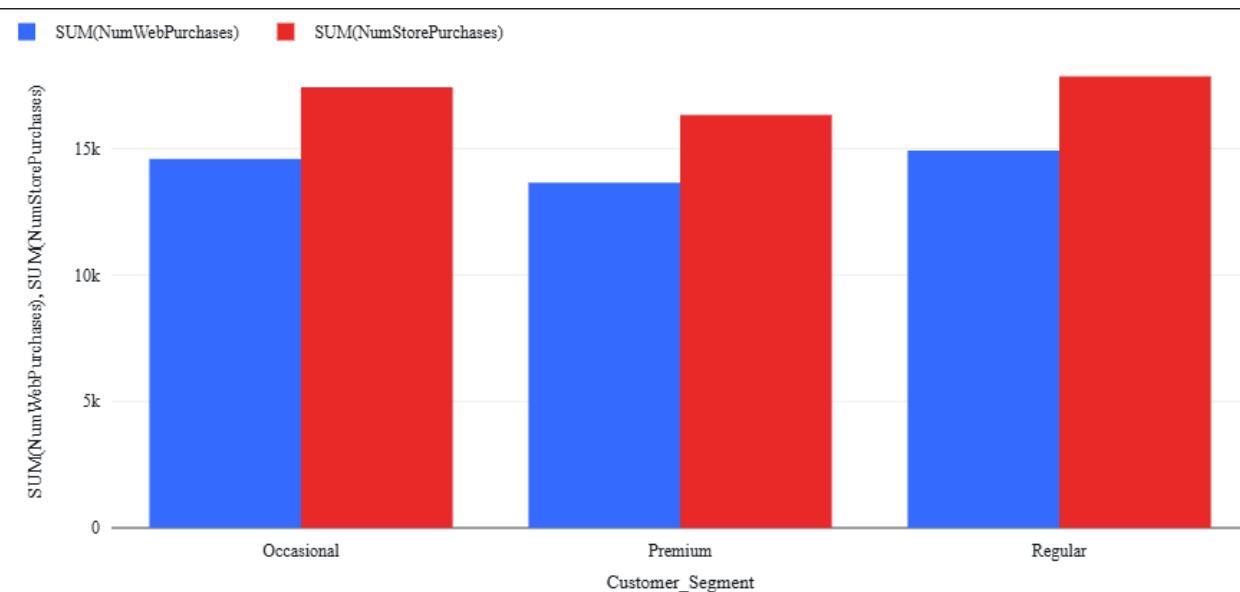
The screenshot shows the Databricks interface for managing jobs and pipelines. On the left, there's a sidebar with various navigation options like Workspace, Recents, Catalog, and Jobs & Pipelines (which is currently selected). The main area displays a job named "job_medallion_etl_full". The "Runs" tab is active, showing a single run named "run_etl_full" which has completed successfully. The "Job details" panel on the right provides information such as Job ID (370658539120945), Creator (Fernando Campos), and Run as (Fernando Campos). It also indicates the lineage (1 upstream table, 2 downstream tables) and performance optimization status. The "Schedules & Triggers" section shows that no triggers are currently defined.

9. Dashboards – Análisis Final

Se generaron 3 visualizaciones desde la tabla Gold:

- ✓ Ventas por Categoría
- ✓ Marketing Spend por Segmento
- ✓ Compras Web vs Tienda





10. Conclusiones Finales

La arquitectura Medallion permite un pipeline organizado, limpio y escalable.

La separación en capas Bronze/Silver/Gold facilitó el mantenimiento y comprensión del flujo.

El uso de Delta Lake garantizó confiabilidad, versionamiento y eficiencia en las tablas.

Unity Catalog brindó gobernanza, control de acceso y orden en los datos.

El workflow permitió automatizar por completo el ETL.

Los dashboards generados demostraron el valor final del pipeline analítico.